

# Part I: Elementary Reinforcement Learning

## 0. Introduction to Reinforcement Learning (RL)

---

### Definitions

- **Markov:**

- a Markov process or markov chain is a memoryless random process, i.e. a sequence of random states  $S_1, S_2, \dots$  with the Markov property.
- a Markov Reward Process (MRP) is a Markov chain with values. It is a tuple  $\langle S, P, R, \gamma \rangle$
- A Markov Decision Process (MDP) is a MRP with decisions. It is an environment where all states are Markov.

Process	Tuple
Markov Process / Chain	$\langle S, P \rangle$
Markov Reward Process (MRP)	$\langle S, P, R, \gamma \rangle$
Markov Decision Process (MDP)	$\langle S, A, P, R, \gamma \rangle$
Partially Observable MDP	$\langle S, A, O, P, R, Z, \gamma \rangle$

with:

- $S$  is a finite set of states
- $A$  is a finite set of actions
- $O$  is a finite set of observations
- $P$  is a state transition probability matrix,  $P_{ss'}^a = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$
- $R$  is a reward function,  $R_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$
- $Z$  is an observation function,  $Z_{s'o}^a = \mathbb{P}[O_{t+1} = o | S_{t+1} = s, A_t = a]$
- $\gamma$  is a discount factor  $\gamma \in [0, 1]$

- **Rewards:**

- a scalar feedback signal  $R_t$ . It indicates how well agent is doing at step t. The agent's job is to maximize cumulative rewards.
- Reward hypothesis: all goals can be described by the maximization of expected cumulative rewards. Reinforcement learning is based on the reward hypothesis.

- **History:** sequence of observations, actions, rewards  $H_t = A_1, O_1, R_1, \dots, A_t, O_t, R_t$

- **Agent** selects the actions and the **environment** selects observations / rewards

- Formally, **the state** is a function of the history -  $S_t = f(H_t)$ . Many states exist:

- environement state  $S_t^e$  : the environment's private representation. Usually not visible to the agent.
- agent's state  $S_t^a$  : the agent's internal representation
- a state  $S_t$  is **Markov** if and only if  $P(S_{t+1}/S_t) = P(S_{t+1}/S_1, S_2, \dots, S_t)$ , i.e. history does not matter

- A **policy** is the agent's behavior, a map from state to action.

- Deterministic policy:  $a = \pi(s)$
- Stochastic policy:  $\pi(a/s) = \mathbb{P}[A_t = a/S_t = s]$  = a probability distribution over actions given states.

- on-policy learning: learn about policy  $\pi$  from experience sampled from policy  $\pi$

- off-policy learning: learn about policy  $\pi$  f from experience sampled from another policy  $\mu$

- $\epsilon$ -greedy methods/policy consist in behaving greedily (picking the action with highest value) most of the time but select randomly with a probability  $\epsilon$  from **among all the actions (including the highest value)** with equal probability (independently of the action-value estimates).

- $\pi(a/s) = \epsilon/m + 1 - \epsilon$  if  $a^* = \underset{a \in A}{\operatorname{argmax}} Q(s, a)$  Note that the probability of picking the highest value is higher than  $1 - \epsilon$  as it can still be picked randomly.
- $\pi(a/s) = \epsilon/m$  otherwise

- a **model** predicts what the environment will do next: transition from state to state and prediction of the next immediate reward. We say that the agent has an internal representation of the environment, called the model.

- Prediction problems: evaluate the future **given a policy**
- Control problems: optimize the future by **finding the best policy**

# 1. Sequential Decision Problems

---

Sequential decision-making problems have the following characteristics:

- **Goal:** select actions to maximise total future rewards
- **Actions:** may have long term consequences
- **Reward:** may be delayed - immediate reward vs. long-term reward trade-off

Two fundamental problems in sequential decision making

- **Reinforcement Learning:**
  - The environment is initially unknown
  - The agent interacts with the environment
  - The agent improves its policy
- **Planning:**
  - A model of the environment is known
  - The agent performs computations with its model (without any external interaction)
  - The agent improves its policy
  - a.k.a. deliberation, reasoning, introspection, pondering, thought, search

A sequential decision problem for a fully observable, stochastic environment with a Markovian transition model and additive rewards is called a **Markov decision process**, or **MDP**, and consists of:

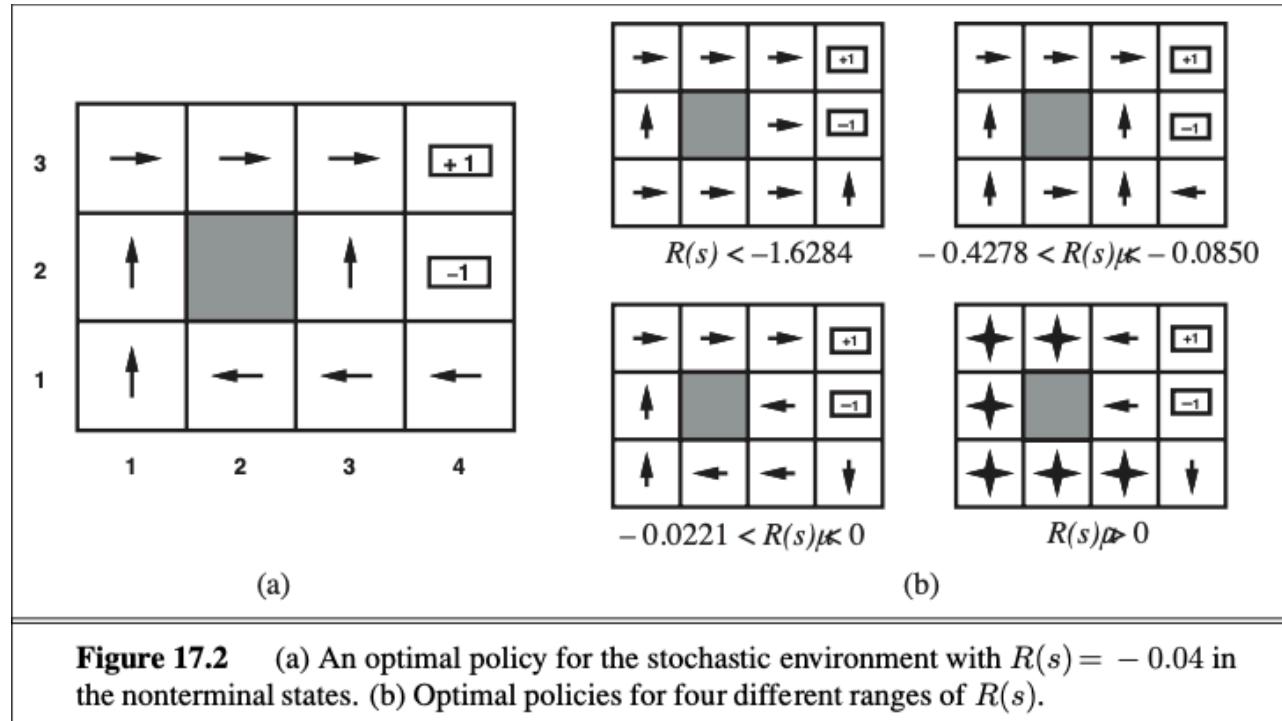
- a set of states (with an initial state  $s_0$ );
- a set ACTIONS(s) of actions in each state A(s), A;
- a transition model  $T(s,a,s') \sim P(s'|s, a)$ ;
- a reward function R(s).
- $\gamma \in [0, 1]$  a discount factor

Therefore, a solution must specify what the agent should do for any state that the agent might reach. **A solution of this kind is called a policy.**

It is traditional to denote a policy by  $\pi$ , and  $\pi(s)$  is the action recommended by the policy  $\pi$  for state s. If the agent has a complete policy, then no matter what the outcome of any action, the agent will always know what to do next.

An optimal policy is a policy that yields the highest expected utility. We use  $\pi^*$  to denote an

optimal policy. Given  $\pi^*$ , the agent decides what to do by consulting its current percept, which tells it the current state  $s$ , and then executing the action  $\pi^*(s)$



**Figure 17.2** (a) An optimal policy for the stochastic environment with  $R(s) = -0.04$  in the nonterminal states. (b) Optimal policies for four different ranges of  $R(s)$ .

**Note:** a policy does not describe a plan but tells us what to do next in a given state. We can still infer a plan based on a policy.

Utilities over time:

- the optimal policy for a finite horizon is non-stationary, i.e. the optimal action in a given state could change over time
- the optimal policy for an infinite horizon is stationary (no reason to behave differently)

We call stationary preferences the fact that, given 2 utilities of sequences  $V_1$  and  $V_2$ ,  $V_1 = \sum_{t=0}^{\infty} U(S_t)$  and  $V_2 = S_0 + \sum_{t=1}^{\infty} U(S'_t)$ : if  $V_1 > V_2$ , then  $\sum_{t=1}^{\infty} U(S_t) > \sum_{t=1}^{\infty} U(S'_t)$

Under stationarity, there are 2 coherent ways to assign utilities to sequences (of states):

- Additive rewards:  $U(s_0, s_1, s_2, \dots) = \sum_{t=0}^{\infty} R(S_t)$
- Discounted rewards  $U(s_0, s_1, s_2, \dots) = \sum_{t=0}^{\infty} \gamma^t R(S_t) \leq \sum_{t=0}^{\infty} \gamma^t R_{\max} = \frac{R_{\max}}{1-\gamma}$

## Fundamental principles / Formulas

### Utility (U), State value (V), State-Action value (Q)

The optimal policy is the policy that maximizes the expected future rewards (by following the

action at each state, returned by the policy)

$$\pi^* = \underset{\pi}{\operatorname{argmax}}(\mathbb{E}[\sum_{t=0}^{\infty} \gamma^t R(S_t)/\pi])$$

The value (or true utility) of a state given a policy is the expected rewards we'll get by following the policy

$$U^\pi(s) = \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t R(s_t)/\pi, s_0 = s]$$

For each state, the optimal policy returns the action that maximizes the weighted value of the next state

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} \sum_{s'} T(s, a, s') U^{\pi^*}(s')$$

The value or true utility of a state is the immediate reward and the discounted future rewards (**Bellman equation**)

$$U(s) = R(s) + \gamma \cdot \max_a \sum_{s'} T(s, a, s') \cdot U(s')$$

It can also be written as:

- a **value** function  $V$  of a state (or **state-value** function) Formally, the state-value function of an MDP is the expected return starting from state  $s$ , and then following policy  $\pi$

$$v_\pi(s) = \mathbb{E}[G_t/S_t = s]$$

- a **quality** function  $Q$  (value of an action) or **action-value** function Formally, the action-value function of an MDP is the expected return starting from state  $s$ , taking action  $a$ , and then following policy  $\pi$

$$q_\pi(s, a) = \mathbb{E}[G_t/S_t = s, A_t = a]$$

- a **continuity** function

# 1. Bellmann Equations (MDPs only)

---

## Bellmann Expectation Equation

### Bellman Expectation Equation

The state-value function can again be decomposed into immediate reward plus discounted value of successor state,

$$v_\pi(s) = \mathbb{E}_\pi [R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s]$$

The action-value function can similarly be decomposed,

$$q_\pi(s, a) = \mathbb{E}_\pi [R_{t+1} + \gamma q_\pi(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a]$$

$$v_\pi(s) = \sum_{a \in A} \pi(a|s) q_\pi(s, a)$$

$$q_\pi(s, a) = R_s^a + \gamma \cdot \sum_{s' \in S} P(s'|s, a) v_\pi(s')$$

By unrolling the equations above, we obtain:

$$v_\pi(s) = \sum_{a \in A} \pi(a|s) \cdot [R_s^a + \gamma \cdot \sum_{s' \in S} P(s'|s, a) v_\pi(s')]$$

$$q_\pi(s, a) = R_s^a + \gamma \cdot \sum_{s' \in S} P(s'|s, a) \cdot \sum_{a' \in A} \pi(a'|s') q_\pi(s', a')$$

## Bellmann Optimality Equation

The optimal state-value function  $v_*(s)$  is the maximum value function over all policies

$$v_*(s) = \max_\pi v_\pi(s)$$

The optimal action-value function  $q_*(s, a)$  is the maximum value function over all policies  
 $q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$

A MDP is "solved" when we know the optimal value function.

## Optimal Policy

### Theorem

*For any Markov Decision Process*

- *There exists an optimal policy  $\pi_*$  that is better than or equal to all other policies,  $\pi_* \geq \pi, \forall \pi$*
- *All optimal policies achieve the optimal value function,  $v_{\pi_*}(s) = v_*(s)$*
- *All optimal policies achieve the optimal action-value function,  $q_{\pi_*}(s, a) = q_*(s, a)$*

Note that there can be more than 1 optimal policy. They all get the same optimal value functions.

An optimal policy can be found by optimizing over  $q_*(s, a)$

$$\pi_*(a/s) = \mathbf{1} \text{ if } a = \operatorname{argmax}_{a \in A} q_*(s, a)$$

$$\pi_*(a/s) = \mathbf{0} \text{ otherwise}$$

There is always a deterministic optimal policy for any MDP.

$$v_*(s) = \max_a q_*(s, a)$$

$$q_*(s, a) = R_s^a + \gamma \cdot \sum_{s' \in S} P(s'/s, a) v_*(s')$$

By unrolling the equations above, we obtain:

$$v_*(s) = \max_a [R_s^a + \gamma \cdot \sum_{s' \in S} P(s'/s, a) v_*(s')]$$

]

$$q_*(s, a) = R_s^a + \gamma \cdot \sum_{s' \in S} P(s'/s, a) \max_{a'} q_*(s', a')$$

Bellman Optimality Equation is non-linear. No closed form solution (in general) Many iterative solution methods exist:

- Value Iteration (DP)
- Policy Iteration (DP)
- Q-learning (off-policy TD control)
- Sarsa (on-policy TD control)

We usually use Bellmann expectation equation to do policy evaluation and the Bellmann optimality equation for control.

## Typology of RL algorithms

From the most to the least supervised and from the least to the most direct learning

- Model-based: from , learns the transition model and calculates  $Q^*$  through a MDP solver and then a policy
- Value-function based or model-free: from , directly learns the  $Q^*$
- policy search: from , directly learns the policy

## 2. Planning by Dynamic Programming (DP)

---

The term dynamic programming (DP) refers to a collection of algorithms that can be used to compute optimal policies given a perfect model of the environment as a Markov decision process (MDP). **Classical DP algorithms are of limited utility in reinforcement learning** both because of their assumption of a **perfect model** and because of their great **computational expense**, but they are still important theoretically. More useful methods attempt to achieve much the same effect as DP, only with less computation and without assuming a perfect model of the environment.

Dynamic programming is a very general solution method for problems with 2 properties:

- Optimal substructure
  - Principle of optimality
  - Optimal solution can be decomposed into subproblems
- Overlapping subproblems
  - Subproblems recur many times
  - Solutions can be cached and reused

MDP satisfy both properties:

- Bellman equation gives recursive decomposition
- Value function stores and reuses solutions

### Policy evaluation

Evaluating a policy consists in the iterative application of Bellmann expectation equation.

Synchronous backups: at each iteration  $k+1$ , for all states, update  $v_{k+1}(s)$  from  $v_k(s')$  where  $s'$  is a successor state of  $s$ .

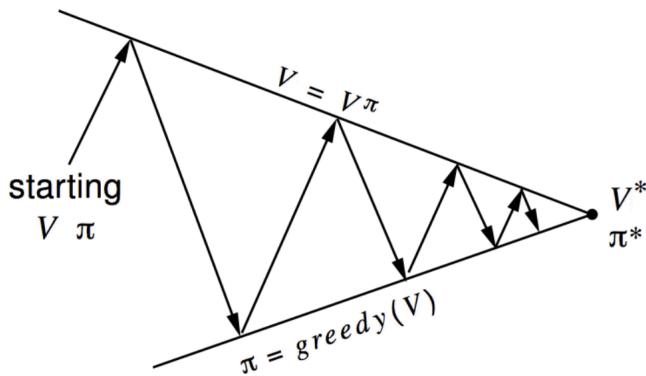
### Policy iteration

Given a policy  $\pi$ ,

- evaluate the policy  $\pi$  (until it converges)
- improve the policy by acting greedily with respect to  $v_\pi$

This process of policy iteration always converges to  $\pi^*$

# Policy Iteration

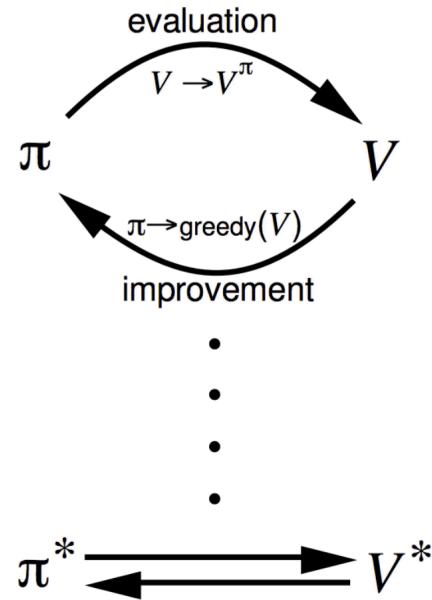


**Policy evaluation** Estimate  $v_\pi$

Iterative policy evaluation

**Policy improvement** Generate  $\pi' \geq \pi$

Greedy policy improvement



The generalised policy iteration principle is that any algorithm can be used for policy evaluation and policy improvement.

Modified policy iteration: stop the iteration before convergence. For example, stop after k iterations of policy evaluation. Value iteration consists in stopping after 1 iteration, i.e 1 iteration of policy evaluation, 1 iteration of policy improvement, 1 iteration of policy evaluation...

## Value iteration

## Deterministic Value Iteration

- If we know the solution to subproblems  $v_*(s')$
- Then solution  $v_*(s)$  can be found by one-step lookahead

$$v_*(s) \leftarrow \max_{a \in \mathcal{A}} \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_*(s')$$

- The idea of value iteration is to apply these updates iteratively
- Intuition: start with final rewards and work backwards
- Still works with loopy, stochastic MDPs

The iteration consists in 2 steps:

- calculating the value function for each action in order to extract the best action.
- selecting the maximum [acting greedily]

Intermediate value functions may not correspond to any policy.

# Synchronous Dynamic Programming Algorithms

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Evaluation
Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

- Algorithms are based on state-value function  $v_\pi(s)$  or  $v_*(s)$
- Complexity  $O(mn^2)$  per iteration, for  $m$  actions and  $n$  states
- Could also apply to action-value function  $q_\pi(s, a)$  or  $q_*(s, a)$
- Complexity  $O(m^2 n^2)$  per iteration

Prediction addresses the question of how good is a given policy. Control optimizes the rewards and finds the optimal policy. Complexity is higher if we use the Q-value function because we move from state-action pair to state-action pair ( $nm \times nm$ ) whereas with V-value, we move from state to state and consider all the actions ( $n \times m \times m$ )

## Asynchronous Dynamic programming

Asynchronous means we update each selected state individually, in any order. If all states continue to be selected, it will converge.

Several methods exist:

- In-place DP (it stores only 1 value function vs. 2 for synchronous)
- Prioritised sweeping (Use magnitude of Bellman error to guide state selection)
- Real-time DP (only states that are relevant to agent)

## Convergence

- Iterative policy evaluation converges to  $v_\pi$  - we also say that  $v_\pi$  converges to a unique fixed point of the Bellmann expectation operator  $T^\pi$

- Iterative policy iteration and iterative value iteration converge to  $v_\pi$

### 3. Model-Free Prediction

---

Dynamic programming is useful to solve a (small) known MDP. Model-Free prediction methods aim at solving unknown MDPs.

#### Monte-Carlo (MC)

- MC methods learn directly from episodes of experience
- MC is model-free: no knowledge of MDP transitions / rewards
- MC learns from complete episodes: no bootstrapping
- MC uses the simplest possible idea: value = mean return
- Caveat: can only apply MC to episodic MDPs (all episodes must terminate)

Monte-Carlo policy evaluation uses empirical mean return instead of expected return.

$V(S) = S(s) / N(s)$  with  $S(s) += G_t$ , the increment total return for state s and  $N(s) += 1$  the increment counter.  $V(s) \rightarrow v_\pi(s)$  as  $N(s) \rightarrow \infty$

- First-visit MC Policy Evaluation: mean return calculated with only 1 visit per episode
- Every-visit MC Policy Evaluation: mean return calculated with several visits per episode

Incremental MC updates consists in tracking a running mean (forget old episodes).

$$\text{Incremental mean : } \mu_k = \frac{1}{k} \sum_{j=1}^k x_j = \mu_{k-1} + \frac{1}{k} \cdot (x_k - \mu_{k-1})$$

In general, `new_estimate<-- old_estimate + step_size [target - old_estimate]`

#### Temporal Difference Learning ( $TD(\lambda)$ )

TD methods learn directly from episodes of experiences. TD learns from incomplete episodes (contrary to MC that requires full episodes), **by bootstrapping**. TD updates a guess towards a guess.

MC has high variance (because of the high variance of the return), zero bias (unbiased estimate of  $v_\pi(s)$ )

- Good convergence properties
- (even with function approximation) Not very sensitive to initial value Very simple to understand and use
- MC does not exploit Markov property (Usually more effective in non-Markov environments)

- MC converges to solution with minimum mean-squared error, i.e. best fit to the observed returns  $\sum_{k=1}^K \sum_{t=1}^{T_k} (G_t^k - V(s_t^k))^2$

TD has low variance (because only dependent on the variance on one random action, transition, reward), some bias (TD target is biased estimate of  $v_\pi(s)$ )

- Usually more efficient than MC
- TD(0) converges to  $v_\pi(s)$
- (but not always with function approximation)
- More sensitive to initial value
- TD exploits Markov property (Usually more efficient in Markov environments)
- TD(0) converges to solution of max likelihood Markov model
  - Solution to the MDP  $\langle S, A, P, R, \gamma \rangle$  that best fits the data

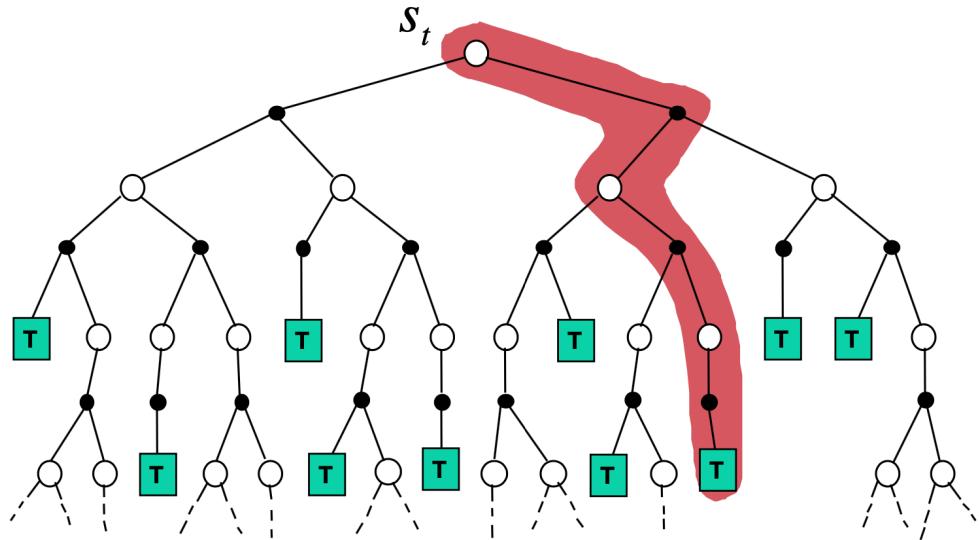
## Backup comparisons

Method	Bootstrapping	Sampling
Dynamic Programming	Yes	No
Monte Carlo	No	Yes
Temporal Difference	Yes	Yes

Visually, bootstrapping = depth of the tree; sampling = width of the tree

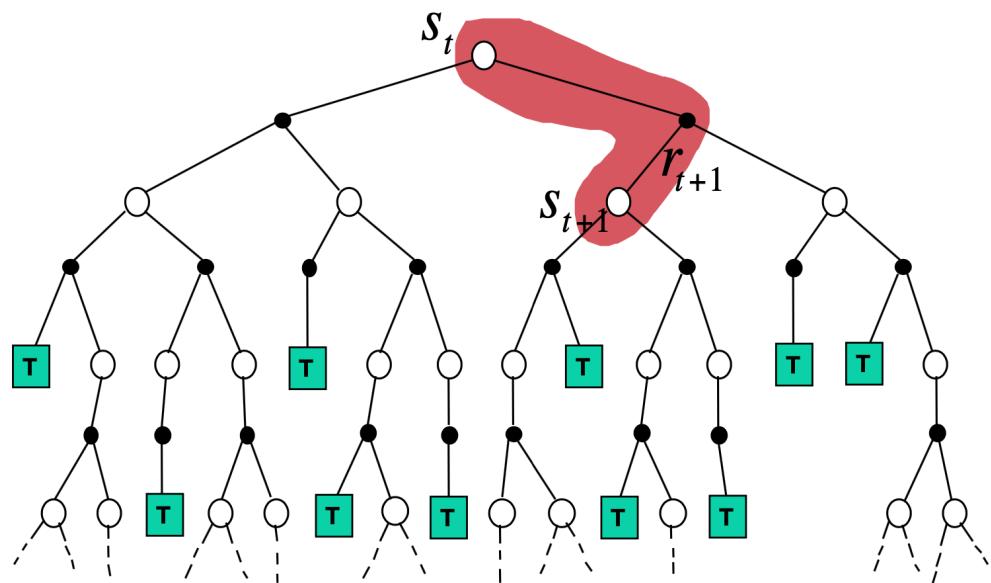
## Monte-Carlo Backup

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$



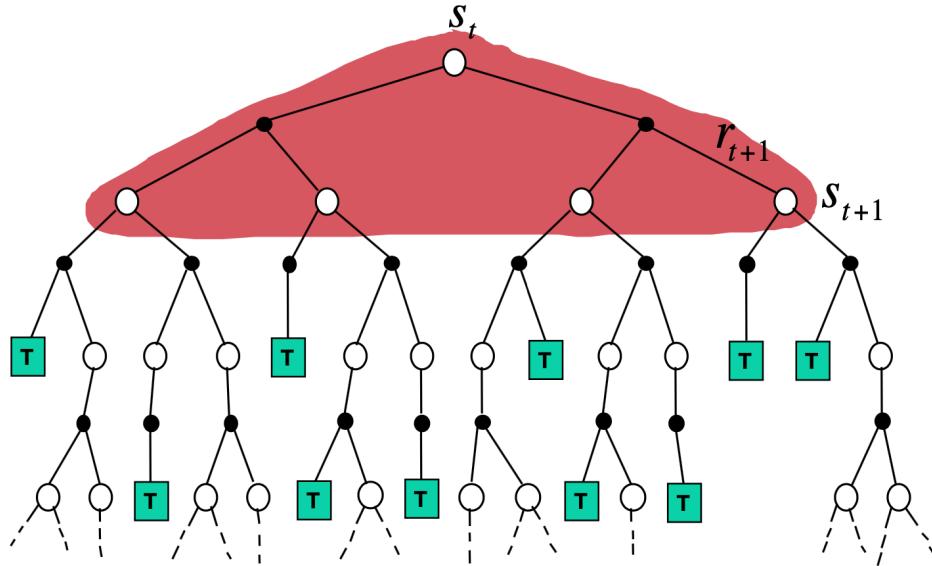
## Temporal-Difference Backup

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$



# Dynamic Programming Backup

$$V(S_t) \leftarrow \mathbb{E}_\pi [R_{t+1} + \gamma V(S_{t+1})]$$



## TD(0) Update Rule

Finds maximum-likelihood estimate (if finite data repeated infinitely often)

**Algorithm:**

Episode T: For all s, at start of episode,  $V_T(s) = V_{T-1}(s)$

For all  $s = s_k$ ,

$$V_T(s_{k-1}) = V_T(s_{k-1}) + \alpha_T \cdot (r_k + \gamma \cdot V_T(s_k) - V_T(s_{k-1}))$$

**then:**  $V_T(s_{k-1}) = \mathbb{E}_{s_k}[r + \gamma \cdot V_T(s_k)]$

## Properties of learning rates

$\lim_{T \rightarrow \infty} V_T(s) = V(s)$  if 2 conditions about  $\alpha_T$  are met (Robbins-Monro sequence of step-sizes):

- $\sum_T \alpha_T = \infty$

- $\sum_T \alpha_T^2 < \infty$

## TD(1) Update Rule

We assume:

- $s_1, s_2, \dots, s_f$  states with reward  $r_k$  between state  $s_k$  and  $s_{k+1}$
- learning rate  $\gamma$
- $e(s)$  is called the eligibility rate

### Algorithm:

Episode T

For all  $s$ ,  $e(s) = 0$  at start of episode,  $V_T(s) = V_{T-1}(s)$

After  $S_{k-1} - > S_k$ , reward  $r_k$ : (Step k)

- $e(s_{k-1}) = e(s_{k-1}) + 1$

For all  $s$ ,

- $V_T(s) = V_T(s) + \alpha_T(r_T + \gamma V_{T-1}(s_T) - V_{T-1}(s_{T-1})). e(s)$
- For all  $k$ :  $e(s) = \gamma \cdot e(s)$

TD(1) is equivalent to outcome-based updates without repeated states with  $V_T(s_k) += \alpha \cdot (\sum_{i=k}^{f-1} \gamma^{i-1} \cdot r_i + \gamma^f \cdot V_{T-1}(S_f) - V_{T-1}(S_k))$

## Generalization: TD( $\lambda$ ) Update rule

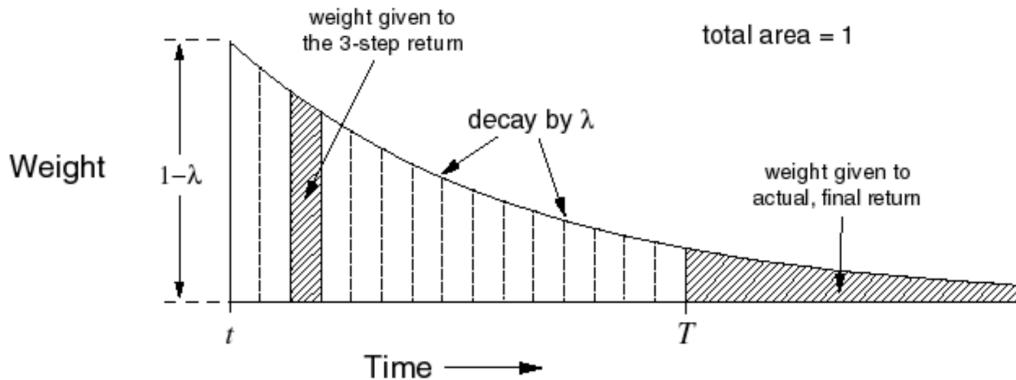
### Forward-view implementation

The principle is to combine different returns with n-step look-ahead,  $G_t^{(n)}$  through a weighted sum. The weight is  $(1 - \lambda) \cdot \lambda^{n-1}$

$$G_t^\lambda = (1 - \lambda) \cdot \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

Forward-view  $TD(\lambda) : V(S_t) \leftarrow V(S_t) + \alpha \cdot (G_t^\lambda - V(S_t))$

## TD( $\lambda$ ) Weighting Function



$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$$

### Backward view implementation (eligibility traces)

Episode T

For all  $s$ ,  $e(s) = 0$  at start of episode,  $V_T(s) = V_{T-1}(s)$

After  $S_{k-1} \rightarrow S_k$ , reward  $r_k$ : (Step k)

- $e(s_{k-1}) = e(s_{k-1}) + 1$

For all  $s$ ,

- $V_T(s) = V_T(s) + \alpha_T(r_T + \gamma V_{T-1}(s_T) - V_{T-1}(s_{T-1})). e(s)$
- For all  $k$ :  $e(s) = \lambda \cdot \gamma \cdot e(s)$

### K-Step Estimators

- $V(s_t) = V(S_t) + \alpha_t. (r_{t+1} + \gamma \cdot V(S_{t+1}) - V(S_t)) = \text{TD}(0)$  [1-step estimator]
- $V(s_t) = V(S_t) + \alpha_t. (r_{t+1} + \gamma \cdot r_{t+2} + \gamma^2 \cdot V(S_{t+2}) - V(S_t)) =$  [2-step estimator]
- ...
- $V(s_t) = V(S_t) + \alpha_t. (r_{t+1} + \gamma \cdot r_{t+2} + \dots + \gamma^{k-1} \cdot r_{t+k} + \gamma^k \cdot V(s_{t+k}) - V(S_t)) =$  [k-step estimator]
- $V(s_t) = V(S_t) + \alpha_t. (r_{t+1} + \gamma \cdot r_{t+2} + \dots + \gamma^{k-1} \cdot r_{t+k} + \dots - V(S_t)) = \text{TD}(1) =$  [infinite-step estimator]

estimator]

## K-Step Estimators and $TD(\lambda)$

$V(S)$  estimated as a weighted combination of step estimators, with weights =  $\lambda^{t-1}(1 - \lambda)$

Empirically, we can observe that error in  $V$  after finite data is:

- highest for  $TD(1)$  but lower for  $TD(0)$
- lowest for  $\lambda$  between 0.3 and 0.7

### Summary of Forward and Backward $TD(\lambda)$

Offline updates	$\lambda = 0$	$\lambda \in (0, 1)$	$\lambda = 1$
Backward view	$TD(0)$ 	$TD(\lambda)$ 	$TD(1)$ 
Forward view	$TD(0)$	Forward $TD(\lambda)$	MC
Online updates	$\lambda = 0$	$\lambda \in (0, 1)$	$\lambda = 1$
Backward view	$TD(0)$ 	$TD(\lambda)$ *	$TD(1)$ *
Forward view	$TD(0)$ 	Forward $TD(\lambda)$ 	MC 
Exact Online	$TD(0)$	Exact Online $TD(\lambda)$	Exact Online $TD(1)$

= here indicates equivalence in total update at end of episode.

## 4. Model-Free Control

---

For most of RL problems, either:

- MDP model is unknown but experience can be sampled
- MDP model is known, but is too big to use, except by samples

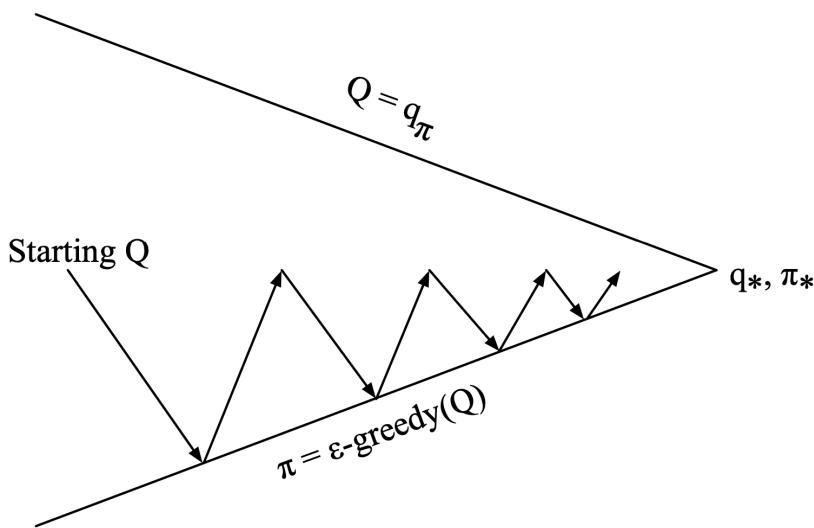
**On-policy learning** describes learning about policy  $\pi$  from experience sampled from  $\pi$ , unlike **off-policy learning** which learns from experience sampled from another policy  $\mu$ .

### On-policy learning

#### Generalised policy iteration (GPI) with Monte-Carlo evaluation: 2 issues:

- We cannot use the greedy policy improvement over  $V(s)$  because it requires a model of MDP  $\pi'(s) = \underset{a \in A}{\operatorname{argmax}} [R_s^a + P_{ss'}^a V(s')]$  ; So we'll use  $Q(s,a)$  instead =>  
$$\pi'(s) = \underset{a \in A}{\operatorname{argmax}} Q(s, a)$$
- We need to maintain some level of exploration. So we'll use a  $\epsilon$ -greedy policy improvement approach.

# Monte-Carlo Control



**Every episode:**

**Policy evaluation** Monte-Carlo policy evaluation,  $Q \approx q_\pi$

**Policy improvement**  $\epsilon$ -greedy policy improvement

The question of this approach is how to balance the need for exploration and the fast convergence into a policy that does not explore anymore. GLIE is an idea to come up with a schedule of exploration to manage this balance (e.g. decaying  $\epsilon$  by  $1/k$ ).

## GLIE Monte-Carlo Control

## Definition

*Greedy in the Limit with Infinite Exploration (GLIE)*

- All state-action pairs are explored infinitely many times,

$$\lim_{k \rightarrow \infty} N_k(s, a) = \infty$$

- The policy converges on a greedy policy,

$$\lim_{k \rightarrow \infty} \pi_k(a|s) = \mathbf{1}(a = \operatorname{argmax}_{a' \in \mathcal{A}} Q_k(s, a'))$$

- For example,  $\epsilon$ -greedy is GLIE if  $\epsilon$  reduces to zero at  $\epsilon_k = \frac{1}{k}$

Here's below one algorithm where we iterate, for every episode, evaluation and improvement - note this is much more efficient than generating thousands of episodes before evaluating the policy (i.e. we can even improve the policy with a single episode):

## GLIE Monte-Carlo Control

- Sample  $k$ th episode using  $\pi$ :  $\{S_1, A_1, R_2, \dots, S_T\} \sim \pi$
- For each state  $S_t$  and action  $A_t$  in the episode,

$$N(S_t, A_t) \leftarrow N(S_t, A_t) + 1$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{1}{N(S_t, A_t)} (G_t - Q(S_t, A_t))$$

- Improve policy based on new action-value function

$$\epsilon \leftarrow 1/k$$

$$\pi \leftarrow \epsilon\text{-greedy}(Q)$$

### Theorem

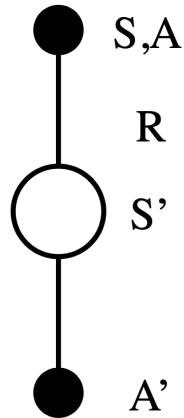
*GLIE Monte-Carlo control converges to the optimal action-value function,  $Q(s, a) \rightarrow q_*(s, a)$*

## TD learning

The main advantage vs. MC is to increase the frequency of the policy improvement - will be done every time-step, instead of at the end of each episode in MC.

The general idea is called SARSA. SARSA name came from the fact that agent takes one step from one state-action value pair to another state-action value pair and along the way collect reward  $R$  (so it's the  $S_t, A_t, R_{t+1}, S_{t+1}$  and  $A_{t+1}$  tuple that creates the term S,A,R,S,A). SARSA is an **on-policy method**. SARSA uses action-value function  $Q$  and follows the policy  $\pi$ . GPI (Generalized Policy Iteration) is used to take action based on policy  $\pi$  ( $\epsilon$ -greedy to ensure exploration as well as greedy to improve the policy).

# Updating Action-Value Functions with Sarsa



$$Q(S, A) \leftarrow Q(S, A) + \alpha (R + \gamma Q(S', A') - Q(S, A))$$

**Update Formula:**

$$Q(S_t, A_t) = Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \cdot Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$

Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

    Initialize  $S$

    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

    Repeat (for each step of episode):

        Take action  $A$ , observe  $R, S'$

        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A'$ ;

    until  $S$  is terminal

**Convergence of SARSA**

## Theorem

*Sarsa converges to the optimal action-value function,  $Q(s, a) \rightarrow q_*(s, a)$ , under the following conditions:*

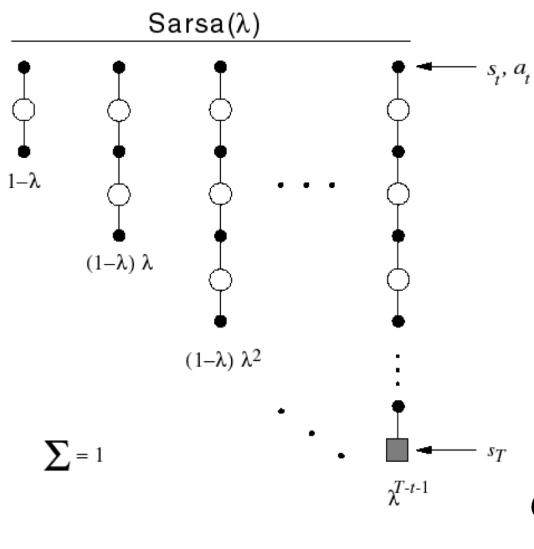
- GLIE sequence of policies  $\pi_t(a|s)$
- Robbins-Monro sequence of step-sizes  $\alpha_t$

$$\sum_{t=1}^{\infty} \alpha_t = \infty$$

$$\sum_{t=1}^{\infty} \alpha_t^2 < \infty$$

## SARSA- $\lambda$ Forward-view

### Forward View Sarsa( $\lambda$ )



- The  $q^\lambda$  return combines all  $n$ -step Q>Returns  $q_t^{(n)}$
- Using weight  $(1 - \lambda)\lambda^{n-1}$

$$q_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} q_t^{(n)}$$

- Forward-view Sarsa( $\lambda$ )

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left( q_t^\lambda - Q(S_t, A_t) \right)$$

## SARSA- $\lambda$ Backward-view

# Sarsa( $\lambda$ ) Algorithm

```
Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
Repeat (for each episode):
     $E(s, a) = 0$ , for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
    Initialize  $S, A$ 
    Repeat (for each step of episode):
        Take action  $A$ , observe  $R, S'$ 
        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)
         $\delta \leftarrow R + \gamma Q(S', A') - Q(S, A)$ 
         $E(S, A) \leftarrow E(S, A) + 1$ 
        For all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ :
             $Q(s, a) \leftarrow Q(s, a) + \alpha \delta E(s, a)$ 
             $E(s, a) \leftarrow \gamma \lambda E(s, a)$ 
         $S \leftarrow S'; A \leftarrow A'$ 
    until  $S$  is terminal
```

## Off-policy learning

Evaluate target policy  $\pi(a/s)$  to compute  $v_\pi(s)$  or  $q_\pi(s, a)$  while following behaviour policy  $\mu(a/s)$ . 2 mechanisms:

- importance sampling:
  - for off-policy MC: extremely high-variance and in practice useless
  - for off-policy TD: works better than MC

## Importance Sampling for Off-Policy TD

- Use TD targets generated from  $\mu$  to evaluate  $\pi$
- Weight TD target  $R + \gamma V(S')$  by importance sampling
- Only need a single importance sampling correction

$$V(S_t) \leftarrow V(S_t) + \alpha \left( \frac{\pi(A_t|S_t)}{\mu(A_t|S_t)} (R_{t+1} + \gamma V(S_{t+1})) - V(S_t) \right)$$

- Much lower variance than Monte-Carlo importance sampling
  - Policies only need to be similar over a single step
- Q-learning: this is the method that works the best

## Q-Learning

- We now consider off-policy learning of action-values  $Q(s, a)$
- **No** importance sampling is required
- Next action is chosen using behaviour policy  $A_{t+1} \sim \mu(\cdot|S_t)$
- But we consider alternative successor action  $A' \sim \pi(\cdot|S_t)$
- And update  $Q(S_t, A_t)$  towards value of alternative action

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (R_{t+1} + \gamma Q(S_{t+1}, A') - Q(S_t, A_t))$$

## Off-Policy Control with Q-Learning

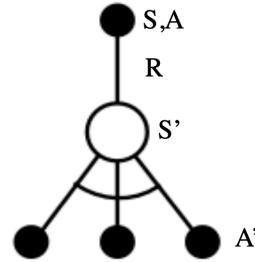
- We now allow both behaviour and target policies to **improve**
- The target policy  $\pi$  is **greedy** w.r.t.  $Q(s, a)$

$$\pi(S_{t+1}) = \operatorname{argmax}_{a'} Q(S_{t+1}, a')$$

- The behaviour policy  $\mu$  is e.g.  **$\epsilon$ -greedy** w.r.t.  $Q(s, a)$
- The Q-learning target then simplifies:

$$\begin{aligned} & R_{t+1} + \gamma Q(S_{t+1}, A') \\ &= R_{t+1} + \gamma Q(S_{t+1}, \operatorname{argmax}_{a'} Q(S_{t+1}, a')) \\ &= R_{t+1} + \max_{a'} \gamma Q(S_{t+1}, a') \end{aligned}$$

## Q-Learning Control Algorithm



$$Q(S, A) \leftarrow Q(S, A) + \alpha \left( R + \gamma \max_{a'} Q(S', a') - Q(S, A) \right)$$

### Theorem

*Q-learning control converges to the optimal action-value function,  
 $Q(s, a) \rightarrow q_*(s, a)$*

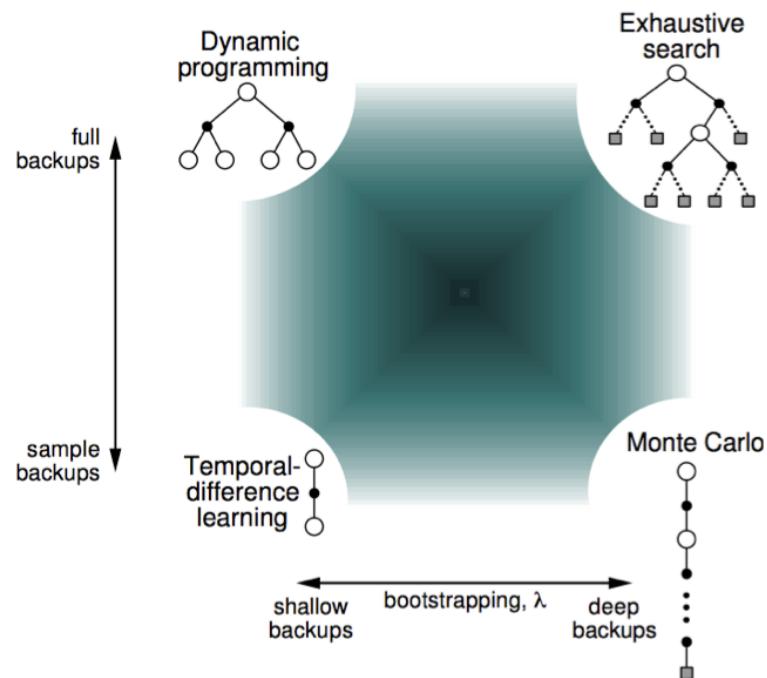
This Q-learning algorithm for control is also called SARSA MAX as it combines SARSA approach and the greedy policy.

There is also the expected SARSA algorithm where we replace the max operation by the expected value of the q\_value of the pair next state, next action.



**Figure 6.4:** The backup diagrams for Q-learning and Expected Sarsa.

# Comparison of RL policy evaluation



## DP and TD

## Relationship Between DP and TD

	<i>Full Backup (DP)</i>	<i>Sample Backup (TD)</i>
Bellman Expectation Equation for $v_\pi(s)$	$v_\pi(s) \leftarrow s$  <b>Iterative Policy Evaluation</b>	 <b>TD Learning</b>
Bellman Expectation Equation for $q_\pi(s, a)$	$q_\pi(s, a) \leftarrow s, a$  <b>Q-Policy Iteration</b>	 <b>Sarsa</b>
Bellman Optimality Equation for $q_*(s, a)$	$q_*(s, a) \leftarrow s, a$  <b>Q-Value Iteration</b>	 <b>Q-Learning</b>

## Relationship Between DP and TD (2)

<i>Full Backup (DP)</i>	<i>Sample Backup (TD)</i>
Iterative Policy Evaluation	TD Learning
$V(s) \leftarrow \mathbb{E}[R + \gamma V(S')   s]$	$V(S) \xleftarrow{\alpha} R + \gamma V(S')$
Q-Policy Iteration	Sarsa
$Q(s, a) \leftarrow \mathbb{E}[R + \gamma Q(S', A')   s, a]$	$Q(S, A) \xleftarrow{\alpha} R + \gamma Q(S', A')$
Q-Value Iteration	Q-Learning
$Q(s, a) \leftarrow \mathbb{E} \left[ R + \gamma \max_{a' \in \mathcal{A}} Q(S', a')   s, a \right]$	$Q(S, A) \xleftarrow{\alpha} R + \gamma \max_{a' \in \mathcal{A}} Q(S', a')$

where  $x \xleftarrow{\alpha} y \equiv x \leftarrow x + \alpha(y - x)$

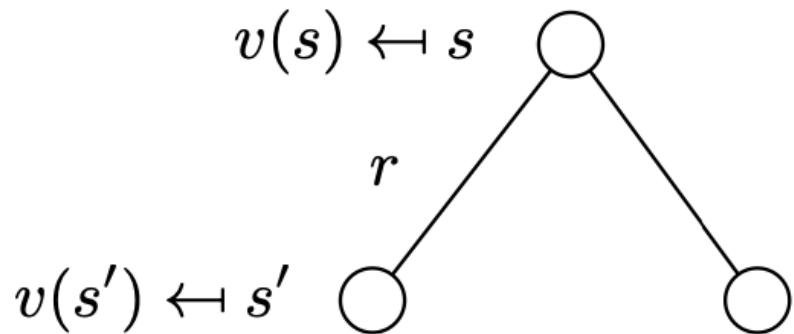
# Appendix

---

## Bellmann Equations for MRPs

- for **MRPs (NO ACTION YET!)**:  $v(s) = \mathbb{E}[G_t/S_t = s] = \mathbb{E}[R_{t+1} + \gamma \cdot v(S_{t+1})/S_t = s]$

$$v(s) = \mathbb{E} [R_{t+1} + \gamma v(S_{t+1}) \mid S_t = s]$$



$$v(s) = \mathcal{R}_s + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'} v(s')$$

- **Matrix Form**

The Bellman equation can be expressed concisely using matrices,

$$v = \mathcal{R} + \gamma \mathcal{P}v$$

where  $v$  is a column vector with one entry per state

$$\begin{bmatrix} v(1) \\ \vdots \\ v(n) \end{bmatrix} = \begin{bmatrix} \mathcal{R}_1 \\ \vdots \\ \mathcal{R}_n \end{bmatrix} + \gamma \begin{bmatrix} \mathcal{P}_{11} & \dots & \mathcal{P}_{1n} \\ \vdots & & \vdots \\ \mathcal{P}_{n1} & \dots & \mathcal{P}_{nn} \end{bmatrix} \begin{bmatrix} v(1) \\ \vdots \\ v(n) \end{bmatrix}$$

So the Bellmann equation is a linear equation that can be solved directly (matrix inversionn)

$$v = R + \gamma P \cdot v$$

$$v = (1 - \gamma \cdot P)^{-1} \cdot R$$

This direct solution is only possible for small MRPs. For large MRPs, there are many iterative methods: Dynamic programming, Monte Carlo, Temporal-Difference learning

## k-bandit problem

- stationary environment:

$$Q_n = \frac{R_1 + R_2 + \dots + R_{n-1}}{n-1}$$

$$Q_{n+1} = \frac{1}{n} \cdot \sum_{i=1}^n R_i = Q_n + \frac{1}{n} [R_n - Q_n]$$

- non-stationary environment:

$$Q_{n+1} = Q_n + \alpha \cdot [R_n - Q_n] = (1 - \alpha)^n Q_1 + \sum_{i=1}^n \alpha (1 - \alpha)^{n-i} R^i$$

It is a weighted average of past rewards and the initial estimate  $Q_1$  with  $\alpha \in (0, 1]$

## Comparison with supervised and unsupervised learning

- Supervised learning is equivalent to finding a function  $f$  based on  $x, y$  pair -  $y = f(x)$ . Also called function approximation

- Unsupervised learning is equivalent to finding a function  $f$  based on  $x$  input data only. Also called clustering description
- Reinforcement learning consists in finding a function, called an optimal policy  $\pi^*$ , based on a set of input data  $(s, a)$  and rewards  $z = r$  for each  $(s, a)$  pair.  $y = f(x) \Leftrightarrow a = \pi^*(s)$ . Main differences:
  - There is no supervisor, only a reward signal.
  - Feedback is delayed, not instantaneous
  - Time really matters
  - Agent takes actions and influences its environment

## Going Further

- Extensions to MDPs
  - Infinite and continuous MDPs
  - Partially Observable Markov Decision Process
  - Belief States (Lecture 2- UCL (David Silver))
  - Ergodic Markov Process (Lecture 2- UCL (David Silver))
  - Average Reward Value Function (Lecture 2- UCL (David Silver))

## References

- repository for RL algorithms:
  - [ShangTong Zhang](#)
  - [Denny Britz](#)
- David Silver courses: <https://www.davidsilver.uk/teaching/>